

MODULE-5

1

INSTRUCTION PIPELINE DESIGN

2

INTRODUCTION

- Instruction execution involves sequence of operations like:-
 - Instruction fetch
 - Decode
 - Operand fetch
 - Execute
 - Write back

PIPELINED INSTRUCTION PROCESSING

○ Fetch stage (F)


- It fetches the instruction from a cache memory
- Ideally one per cycle

○ Decode stage (D)


- It reveals the instruction function to be performed
- Identifies the resources needed
- The resources are:-
 - General purpose registers
 - Buses
 - Functional units

○ Issue stage (I)

- It reserves the resources
- The operands are read from the registers during the issue stage



(a) A seven-stage instruction pipeline




(b) In-order instruction issuing

○ Execute stages (E)


- Instructions are executed in one or several execute stages

○ Write back stage (W)

- This stage is used to write the results into the registers
- Memory load or store operations are treated as part of execution



(a) A seven-stage instruction pipeline



(b) In-order instruction issuing

- Aim
 - $X = y + z$
 - $A = b * c$
- Shaded boxes corresponds to idle cycles
 - This situation happens when instruction issues are blocked due to
 - Resource latency or
 - Due to conflicts of data dependency
- First 2 load instructions issue on consecutive cycles
- Add is dependent on both loads
 - Hence it must wait 3 cycles before the data Y and Z is loaded


TOTAL TIME REQUIRED

- It is measured beginning at cycle 4, when the first instruction starts execution until cycle 20 when the last instruction starts execution
- This timing measure eliminates the undue effects of pipeline startup or draining delays
- Total time= 17 cycles

INSTRUCTION REORDERING

- To improve the timing, **instruction issuing order** is changed
- This eliminates unnecessary delays due to dependence
- All the load operations are issued first
- This reduces the timing to **11 cycles**

$R1 \leftarrow \text{Mem}(Y)$
 $R2 \leftarrow \text{Mem}(Z)$
 $R4 \leftarrow \text{Mem}(B)$
 $R5 \leftarrow \text{Mem}(C)$
 $R3 \leftarrow (R1) + (R2)$
 $R6 \leftarrow (R4) * (R5)$
 $\text{Mem}(x) \leftarrow (R3)$
 $\text{Mem}(A) \leftarrow (R6)$



(c) Reordered instruction issuing

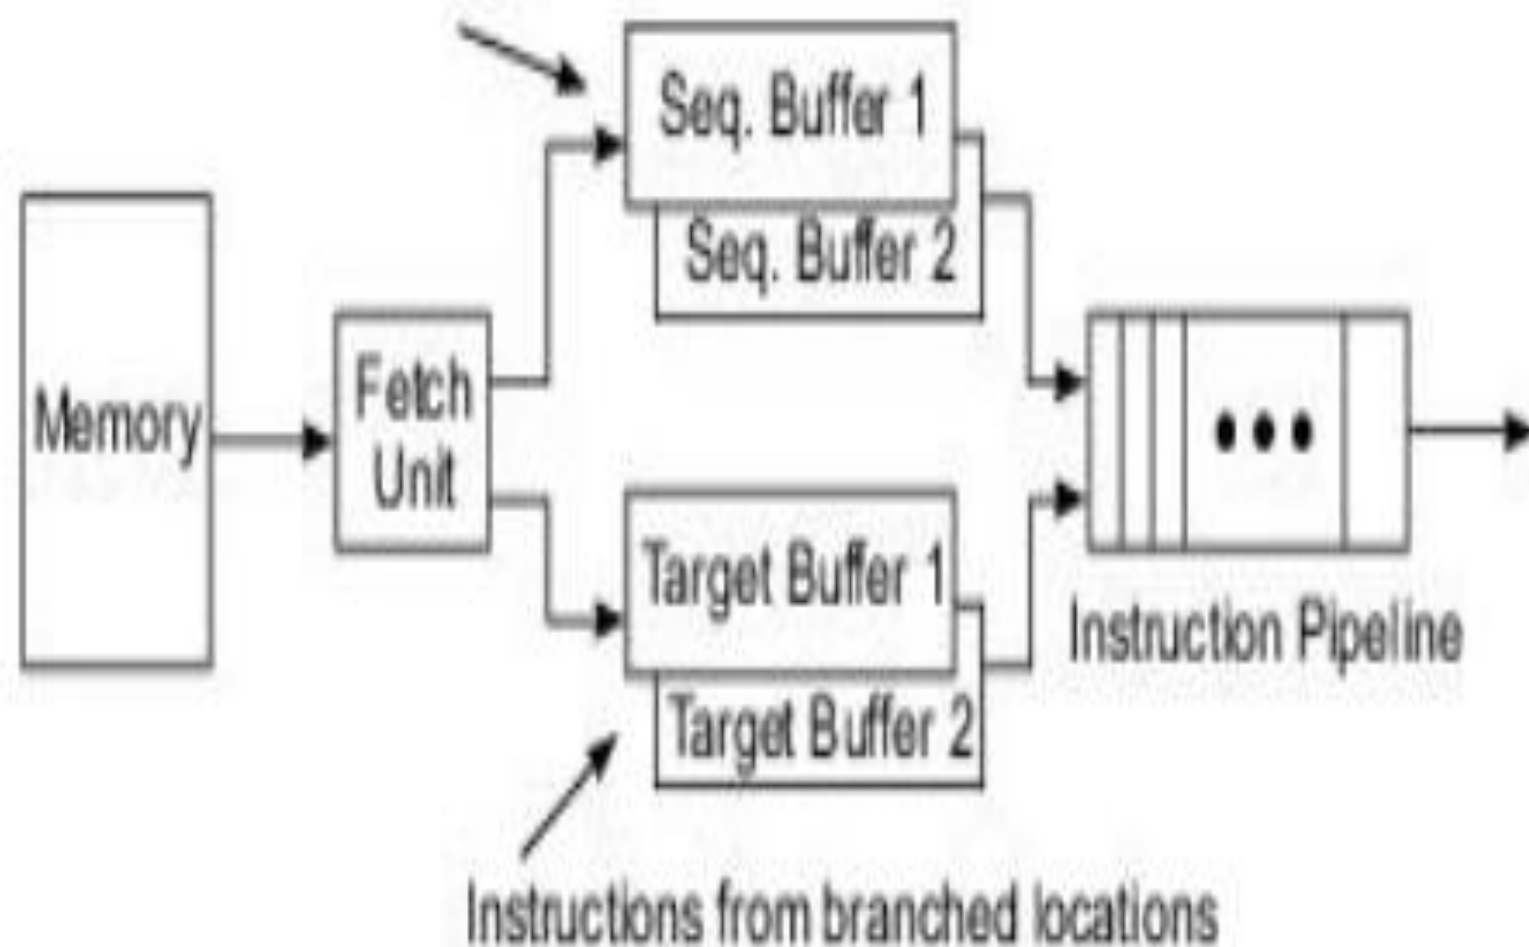
MECHANISM FOR INSTRUCTION PIPELINING

- Techniques are used to **smoothen** the pipeline flow and to **reduce the bottlenecks** of unnecessary memory accesses
 - Pre-fetch buffers
 - Multiple functional units
 - Internal data forwarding
 - Hazard avoidance

PRE-FETCH BUFFERS

- A block of consecutive instructions are fetched into a pre-fetch buffer in one memory access time
- 3 types of buffers are used to match the instruction fetch rate to the pipeline consumption rate
 - Sequential buffer
 - Target buffer
 - Loop buffer

Sequential instructions indicated by program counter



○ Sequential buffer

- Sequential instructions are loaded into a pair of sequential buffers for in-sequence pipelining

○ Target buffer

- Instructions from a branch target are loaded into a pair of target buffers for out of sequence pipelining

○ Both buffers operate in FIFO fashion

○ Buffers alternate to avoid collision

○ Buffers become a part of pipeline as additional stages

○ Conditional branch instruction

- This causes both sequential buffer & target buffer to fill with instructions
- After the branch condition is checked, appropriate instructions are taken from one of the 2 buffers
- Instruction in the other buffer is discarded

○ Loop buffer

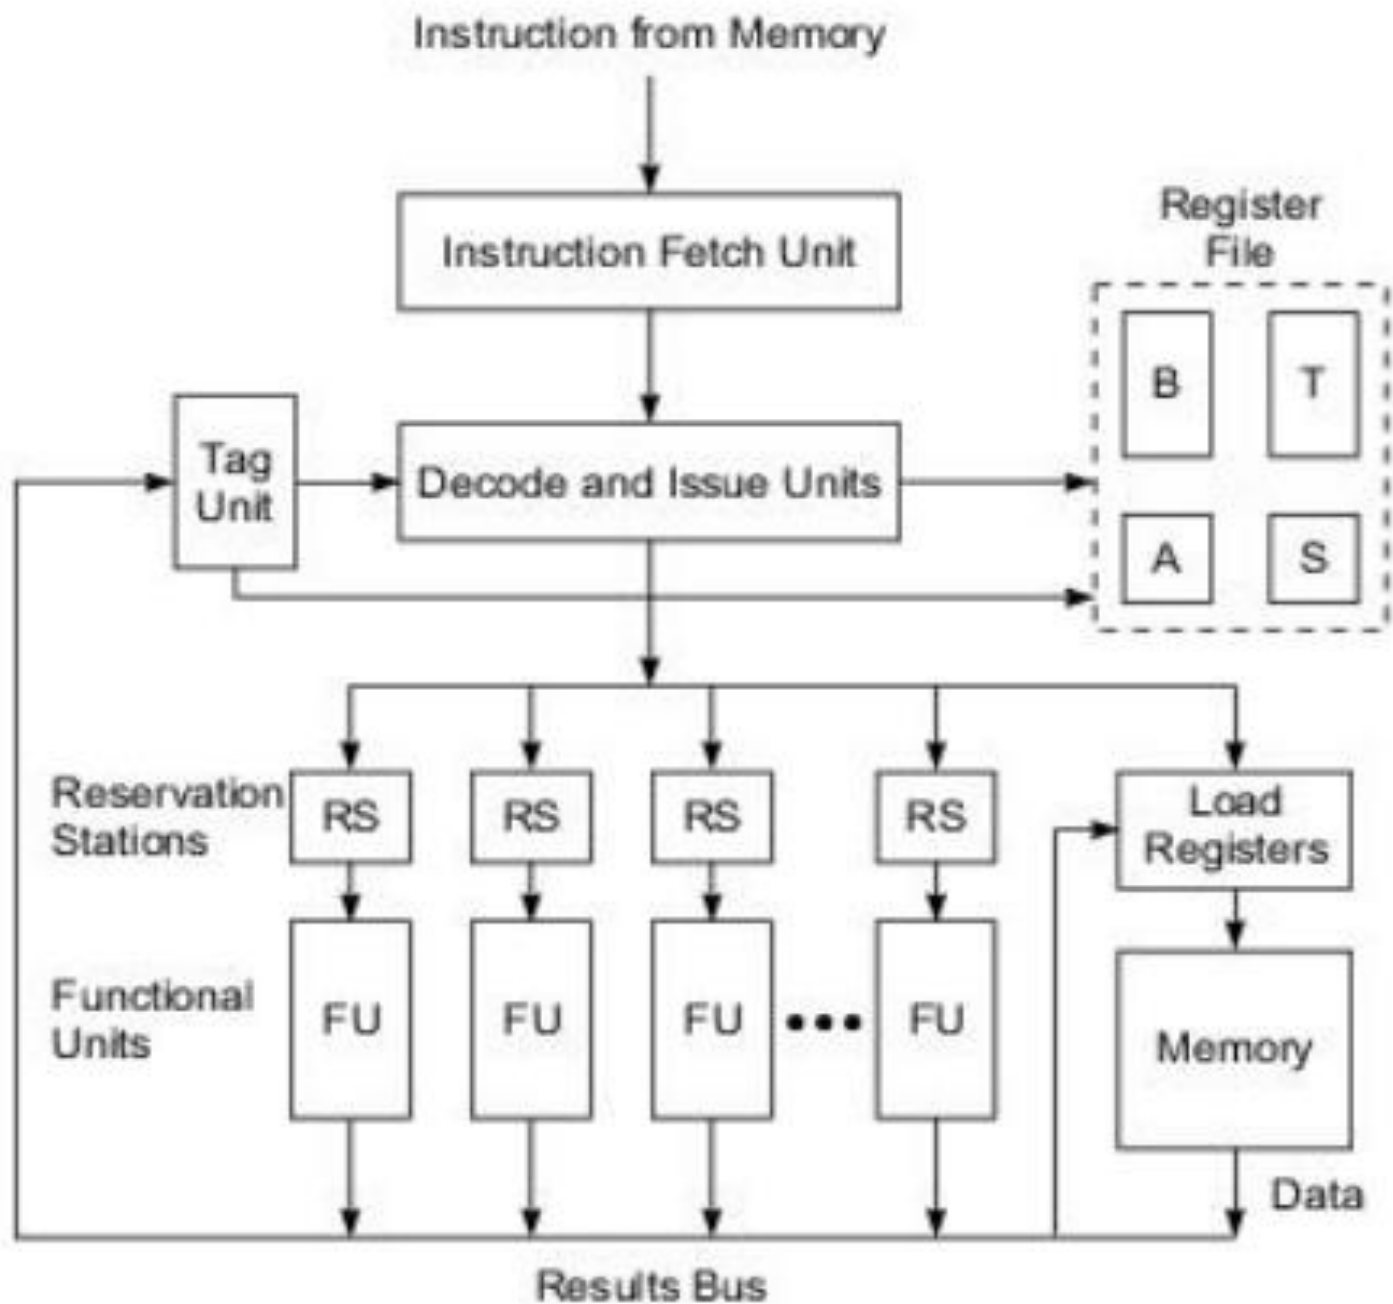
- This buffer holds the sequential instructions contained in a small loop
- Prefetched instructions in the loop body will be executed repeatedly until all iterations complete executions

M U L T I P L E F U N C T I O N A L U N I T S

- Sometimes a particular pipeline stage becomes a bottleneck
 - This may happen to the row which contain max no: of checkmarks in a reservation table
- This bottleneck problem can be alleviated by using multiple copies of same stage simultaneously
- This leads to the use of multiple execution units in a pipelined processor design

PIPELINED PROCESSOR WITH MULTIPLE FUNCTIONAL UNITS

- This architecture consist of a Reservation Station along with each functional unit
 - They are used to resolve the data or resource dependency among successive instructions entering the pipeline
- Operations wait in the RS until their data dependency are resolved
- RS also serve as a buffer to interface the pipeline functional units with decode & issue unit
- RS is uniquely identified using a tag




- Tag is monitored by a tag unit
 - They check the tags of all currently used registers or RS
 - This allows to resolve conflicts between source & destination registers assigned for multiple instructions
- When dependencies are resolved, multiple FU execute in parallel
- This reduces the bottleneck in the execution stages of instruction pipeline

INTERNAL DATA FORWARDING


- Throughput of pipelined processor is improves using internal data forwarding among multiple FU
- **A store-load forwarding**, which consist of a load operation (LD R2,M) can be replaced by a move operation from R1 to R2
 - MOVE R2,R1
- Register transfer is faster than memory access
- Hence data forwarding will reduce memory traffic
- This results in short execution time



(a) Store-load forwarding



(b) Load-load forwarding



- In load-load forwarding, second load is replaced with move operation
 - LD R2,M → MOVE R2,R1

HAZARD AVOIDANCE

- Out of order execution of instructions may lead to certain hazards
- Reason is the read & write operation on shared variables by different instruction in pipeline
- I and J are 2 instructions
 - J follows I logically according to program order
 - If the execution order of these instruction violates the program order, incorrect results may be read or written
 - This leads to hazards

TYPES OF HAZARDS

- 3 types of logical hazards
 - Read after write (RAW)
 - Write after read (WAR)
 - Write after write (WAW)
- Consider 2 instructions I and J
- J follows I in program order
- $D(I)$ (Domain of I)
 - contains the i/p set to be used by the instruction I
- $R(I)$ (range of I)
 - o/p set of instruction I

READ AFTER WRITE HAZARD

- RAW **hazard** refers to a situation where an instruction tries to read a result that has not yet been calculated or retrieved
- J tries to read a source, before I writes to it
- RAW corresponds to flow dependence
- Eg:
 - I: $R2 \leftarrow R1 + R3$
 - J: $R4 \leftarrow R2 + R3$

WRITE AFTER READ **HAZARD** (WAR)

- This hazard occurs when instruction J tries to write to a destination, before it is read by I
- WAR corresponds to antidependence
- Eg:
 - I: $R4 \leftarrow R1 + R5$
 - J: $R5 \leftarrow R1 + R2$

WRITE AFTER WRITE HAZARD

- J tries to write an operand before it is written by I
- WAW corresponds to output dependence
- Eg:
 - I: $R2 \leftarrow R4 + R7$
 - J: $R2 \leftarrow R1 + R3$


- Hazards must be prevented before instructions enter the pipeline
- This is done by holding instruction J, until the dependency on I is resolved

DYNAMIC INSTRUCTION SCHEDULING

30

TOMASULO'S ALGORITHM

- It is a hardware scheme for dependence resolution
- It was first implemented in floating point units of IBM360/91 processor
- In model91 processor,
 - 3 Reservation stations were used in a floating point adder
 - 2 pairs of RS were used in floating point multiplier
- This scheme resolved data dependencies and resource conflicts using register tagging
 - This helped to allocate or deallocate source & destination registers
- This scheme was applied to processors having few floating point registers



WORKING

- An issued instructions whose operands are not available is forwarded to RS associated with the FU it will use
 - It waits until the data dependencies are resolved & the operands become available
- When all operands of an instruction are available,
 - instruction is dispatched to the functional unit for execution.
- All working registers are tagged.
 - If a source register is busy when an instruction reaches the issue stage, the tag for the source register is forwarded to an RS.
 - When register data becomes available, it also reaches the RS which has the same tag.

- The dependence is resolved by monitoring the result bus
 - When an instruction has completed execution, the result along with its tag appears on the result bus.
 - The registers as well as the RS' s monitor the result bus and update their contents when a matching tag is found.

EXAMPLE


○ Aim

- $X=Y+Z$
- $A=B*C$

total execution time= 13 cycles

$R1 \leftarrow \text{Mem}(Y)$
 $R2 \leftarrow \text{Mem}(Z)$
 $R3 \leftarrow (R1)+(R2)$
 $\text{Mem}(x) \leftarrow (R3)$
 $R1 \leftarrow \text{Mem}(B)$
 $R2 \leftarrow \text{Mem}(C)$
 $R3 \leftarrow (R1)*(R2)$
 $\text{Mem}(A) \leftarrow (R3)$

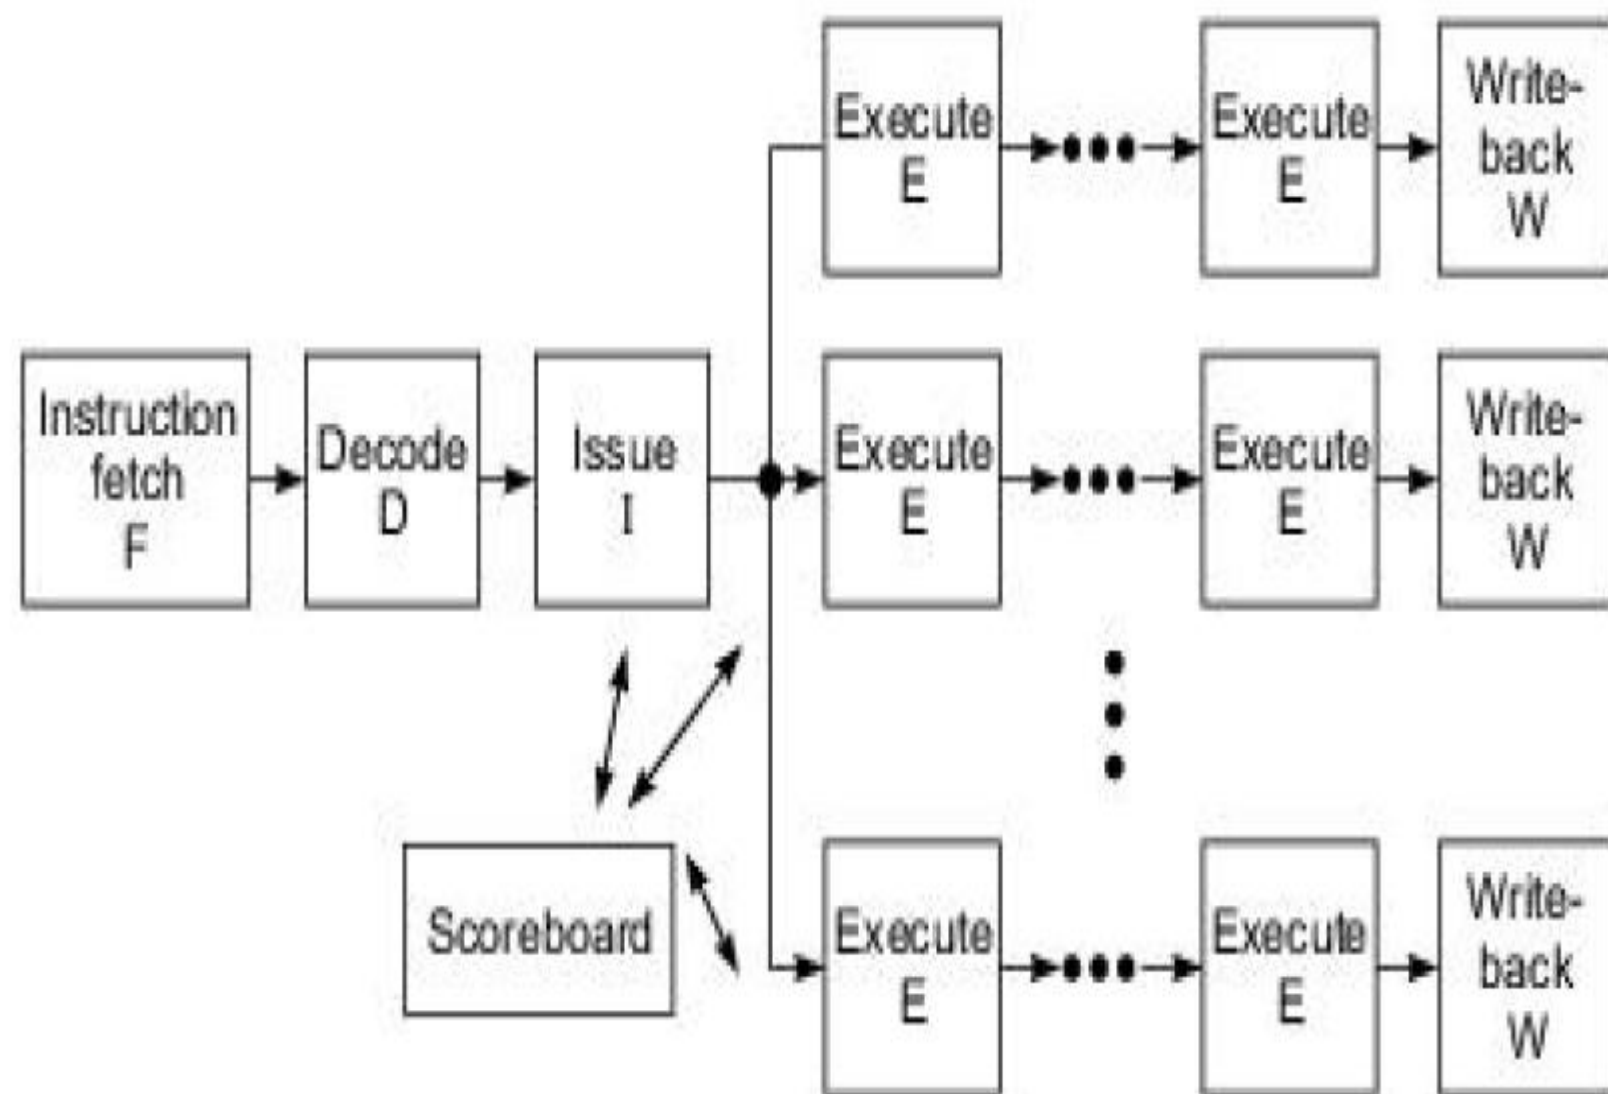
(a) Minimum-register machine code



(b) The pipeline schedule

CDC SCORE BOARDING

- CDC6600 is a high-performance computer that used dynamic instruction scheduling hardware
- In this model, multiple functional units appeared as multiple execution pipelines
- The processor had instruction buffers for each execution unit.
- Instructions were issued to available functional units regardless of whether register input data was available



(a) A CDC 6600-like processor

- The instruction waits in a buffer until its data is produced by other instructions.
- To control the correct routing of data between execution units and registers, there is a centralized control unit known as the **scoreboard**.


SCOREBOARD

- The scoreboard is a centralized control logic which keeps track of the status of registers and multiple functional units.
 - It kept track of the registers needed by instructions waiting for the various functional units.
 - When all registers had valid data, the scoreboard enabled the instruction execution.
 - Similarly, when a functional unit finished execution, it signaled the scoreboard to release the resources.

EXAMPLE

- Aim
 - $X=Y+Z$
 - $A=B*C$
- The add instruction is issued to its functional unit before its registers are ready.
- It then waits for its input register operands.
- The scoreboard routes the register values to the adder unit when they become available.
- In the meantime, the issue stage is not blocked, so other instructions can bypass the blocked add.

$R1 \leftarrow \text{Mem}(Y)$
 $R2 \leftarrow \text{Mem}(Z)$
 $R3 \leftarrow (R1) + (R2)$
 $\text{Mem}(x) \leftarrow (R3)$
 $R4 \leftarrow \text{Mem}(B)$
 $R5 \leftarrow \text{Mem}(C)$
 $R6 \leftarrow (R4) * (R5)$
 $\text{Mem}(A) \leftarrow (R6)$



CONCLUSION

- Dynamic instruction scheduling was implemented only in high-end mainframes or supercomputers in the past.
- Most microprocessors used static scheduling.
- But the trend has changed over the last two decades.
- Today , RI SC and superscalar processors are built with hardware support of dynamic scheduling at runtime.

BRANCH HANDLING TECHNIQUES

- The performance of pipelined processors is limited by data dependences and branch instructions
- Effects of branching
 - 3 terms used to analyze the effects of branching are:-
 - Branch taken
 - Branch target
 - Delay slot

○ Branch taken

- The action of fetching a non sequential or remote instruction after branch instruction is called branch taken



○ Branch target

- The instruction to be executed after a branch taken is called a branch target.

○ Delay slot

- It is the no: of pipeline cycles wasted between a branch taken and the fetching of its branch target
- It is denoted by b .
- Generally $0 \leq b \leq k-1$, where k is the no: of pipeline stages

- When a branch is taken,
 - all the instructions following the branch in the pipeline become useless
 - Hence it will be drained from the pipeline.
 - This implies that a branch taken causes the pipeline to be flushed.
 - This lead to wastage of a no: of useful cycles.
- Branch taken causes I_{b+1} through I_{b+k-1} to be drained from the pipeline.



TOTAL EXECUTION TIME

- Total execution time of n instruction including the effect of branching
 - $T_{\text{eff}} = k\tau + (n-1)\tau + pqnb\tau$
- $pqnb\tau \rightarrow$ branch penalty
- $p \rightarrow$ probability of a conditional branch instruction in a typical instruction stream
- $q \rightarrow$ probability of a successfully executed conditional branch instruction
- $b \rightarrow$ delay slot size
- $k \rightarrow$ no: of pipeline stages
- $n \rightarrow$ no: of instructions
- $\tau \rightarrow$ clock cycle

EFFECTIVE PIPELINE THROUGHPUT

- Effective pipeline throughput including the influence of branching is:-

- $$\text{Heff} = \frac{n}{T_{\text{eff}}} = \frac{nf}{k+n-1+pqnb}$$

BRANCH PREDICTION

- A branch can be predicted using 2 methods
 - Static branch prediction → done by compiler
 - Dynamic branch prediction
- In static branch prediction,

STATIC BRANCH PREDICTION

- Branch is predicted based on branch instruction type
- Static strategy requires the collection of
 - frequency & probabilities of branch taken
 - Branch types across program
- Static branch prediction not very accurate

DYNAMIC BRANCH PREDICTION

- In dynamic branch prediction, branch is predicted based on the branch history
- Better than static prediction
 - It uses recent branch history to predict whether the branch will be taken next time or not
 - It also specifies when a branch occurs
- To predict this:-
 - Use the entire history of branch
- This is infeasible to implement
 - Hence dynamic prediction is determined with limited recent history

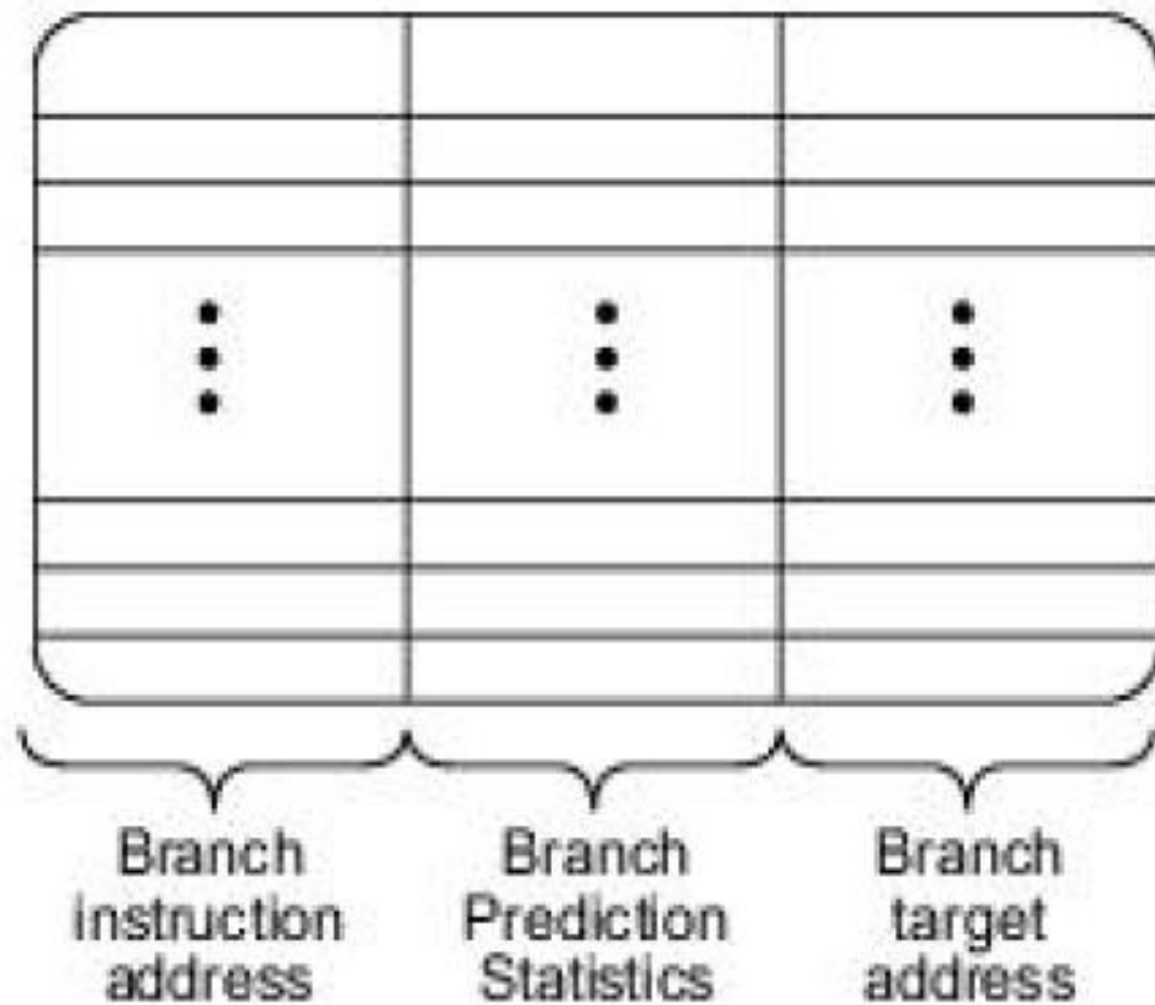
CLASSES OF DYNAMIC BRANCH STRATEGIES

- Class1
 - Predicts the branch direction based on information found at the decode stage
- Class 2
 - Predicts the branch direction at the stage when effective address of the branch target is computed
 - It uses a cache to store the target addresses
- Class 3
 - Uses a cache to store the target instructions at fetch stage

- Dynamic prediction requires additional hardware
 - This h/w keep track of the past behavior of the branch instructions at run time.
 - The amount of history recorded should be relatively small.
 - Otherwise, the prediction logic becomes too costly to implement.

BRANCH TARGET BUFFER

- Lee and Smith suggested the use of a branch target buffer, to implement branch prediction
- It is used to hold recent branch information
- The BTB entry contains the information which will guide the prediction.
- It includes the following:-
 - address of the branch target
 - Address of branch instruction
 - Branch prediction statistics
- Prediction information is updated upon completion of the current branch
- BTB can be extended to store not only the branch target address but also the target instruction itself
 - This is to allow zero delay in converting conditional branches to unconditional branches.



(a) Branch target buffer organization



ARITHMETIC **PIPELINE** DESIGN

56

COMPUTER ARITHMETIC PRINCIPLES

- Pipelining techniques can be applied to speed up numerical arithmetic computations
- In computers arithmetic operations are performed with finite precision
 - Finite precision implies that numbers exceeding the limit must be truncated or rounded
 - This is to provide a precision within the no: of significant bits allowed.
- Fixed-point arithmetic operates on a fixed range of numbers
- Floating-point arithmetic operates over a dynamic range of numbers.


FIXED POINT OPERATIONS

- They are represented internally in machines as
 - sign-magnitude,
 - 1's complement
 - 2's complement notation
- Most computers use the two's complement notation because of its unique representation of all numbers (including zero).
- Four primitive arithmetic operations are:-
 - Add
 - Subtract
 - Multiply
 - Divide

FLOATING POINT OPERATIONS

- Floating -point number X is represented by a pair (m,e)
 - $m \rightarrow$ mantissa or fraction
 - $e \rightarrow$ exponent with an implied base or radix
- Floating point numbers, exceeding the exponent range leads to error conditions, called overflow or underflow
- IEEE has developed standard formats for 32- and 64-bit floating numbers known as the IEEE 754 standard

IEEE 754 STANDARD



STATIC ARITHMETIC PIPELINES

- Static arithmetic pipelines are designed to perform a fixed function
- They are hence called as unifunctional pipelines
- Most of arithmetic pipelines are designed to perform fixed functions.
- ALU perform fixed-point and floating-point operations separately.
 - The fixed-point unit is also called as integer unit
 - The floating-point unit can be built either as part of the central processor or on a separate coprocessor.
- Pipelining in scalar arithmetic pipelines are controlled by software loops

- Scalar and vector arithmetic pipelines differ in the areas of
 - register files
 - control mechanisms
- Vector hardware pipelines are built as
 - add-on options to a scalar processor
 - attached processor driven by a control processor

ARITHMETIC PIPELINE STAGES

- Depending on the function to be implemented, pipeline stages in an arithmetic unit require different hardware logic.
- Arithmetic operations like
 - Add
 - Subtract
 - multiply
 - Divide
 - Squaring
 - square root etc
- are implemented with the basic add and Shifting operations
- Hence the core arithmetic stages require some form of hardware to perform add and shift.

PIPELINE STAGES OF A 3 STAGE FLOATING POINT ADDER


- Stage 1: exponent comparison and equalization
 - It is implemented with an integer adder and some shifting logic
- Stage 2: fraction addition
 - Implemented using a high-speed carry lookahead adder
- Stage 3: Fraction normalization and exponent readjustment
 - Implemented using a shifter and another addition logic

CPA AND CSA

- Arithmetic or logical shifts can be easily implemented with shift registers.
- High-speed addition requires either the use of a
 - Carry-propagation adder (CPA) or
 - Carry save adder (CSA)
- CPA adds two numbers and produces an arithmetic sum
 - In a CPA, the carries generated in successive digits are allowed to propagate from the low end to the high end,
 - Propagation is done using
 - ripple carry propagation
 - carry look ahead technique.

e.g. $n=4$

$$\begin{array}{r} A = 1\ 0\ 1\ 1 \\ +) B = 0\ 1\ 1\ 1 \\ \hline S = 1\ 0\ 0\ 1\ 0 = A + B \end{array}$$



(a) An n -bit carry-propagate adder (CPA) which allows either carry propagation or applies the carry-lookahead technique


- CSA add three input numbers and produce one sum output and a carry output
 - In a CSA, the carries are not allowed to propagate but instead are saved in a carry vector
 - An n bit CSA is specified as:-
 - Let X, Y, and Z be three n-bit input numbers, expressed as $X=(x_{n-1}, x_{n-2}, \dots, x_1, x_0)$
 - The CSA performs bitwise operations simultaneously on all columns of digits
 - It produces two n-bit output numbers, denoted as S^b and C
 - $S^b \rightarrow$ bitwise sum
- $S^b=(0, S_{n-1}, S_{n-2}, \dots, S_1, S_0)$
- $C=(C_n, C_{n-1}, \dots, C_1, C_0)$
 - Leading bit of S^b is always 0
 - Trail bit of C is always 0

○ The i/p o/p relationship is expressed as

- $S_i = x_i \text{ XOR } y_i \text{ XOR } z_i$
- $C_{i+1} = x_i y_i \text{ OR } y_i z_i \text{ OR } z_i x_i$

e.g. $n=4$

$$\begin{array}{r}
 X = \quad 0 \ 0 \ 1 \ 0 \ 1 \ 1 \\
 Y = \quad 0 \ 1 \ 0 \ 1 \ 0 \ 1 \\
 \oplus Z = \quad 1 \ 1 \ 1 \ 1 \ 0 \ 1 \\
 \hline
 S^b = 0 \ 1 \ 0 \ 0 \ 0 \ 1 \ 1 \\
 +) C = 0 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0 \\
 \hline
 S = 1 \ 0 \ 1 \ 1 \ 1 \ 0 \ 1 = S^b + C = X + Y + Z
 \end{array}$$




MULTIPLY PIPELINE DESIGN

- Consider following operation
 - $A * B = P$
- A and B are 8 bit inputs
- $P \rightarrow$ 16 bit product
- Fixed point multiplication is the summation of 8 partial products
 - $P = P_0 + P_1 + P_2 + \dots + P_7$

									1	0	1	1	0	1	0	1	=	A
								×)	1	0	0	1	0	0	1	1	=	B
<hr/>																		
									1	0	1	1	0	1	0	1	=	P_0
								1	0	1	1	0	1	0	1	0	=	P_1
							0	0	0	0	0	0	0	0	0	0	=	P_2
					0	0	0	0	0	0	0	0	0	0	0	0	=	P_3
				1	0	1	1	0	1	0	1	0	0	0	0	0	=	P_4
			0	0	0	0	0	0	0	0	0	0	0	0	0	0	=	P_5
		0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	=	P_6
+) 1	0	1	1	0	1	0	1	0	1	0	0	0	0	0	0	0	=	P_7
<hr/>																		
0	1	1	0	0	1	1	1	1	1	1	0	1	1	1	1	1	=	P

PIPELINE UNIT FOR FIXED POINT MULTIPLICATION

- Stage S1
 - generates all eight partial products, ranging from 8 bits to 15 bits simultaneously.
- Stage S2
 - It is made of two levels of four CSA
 - It merges eight numbers into four numbers
- Stage S3
 - consists of two CSAs,
 - It merges four numbers from S2 into two 16-bit numbers
- Stage S4
 - It is a CPA, which adds up the last two numbers to produce the final product P.



Captions:

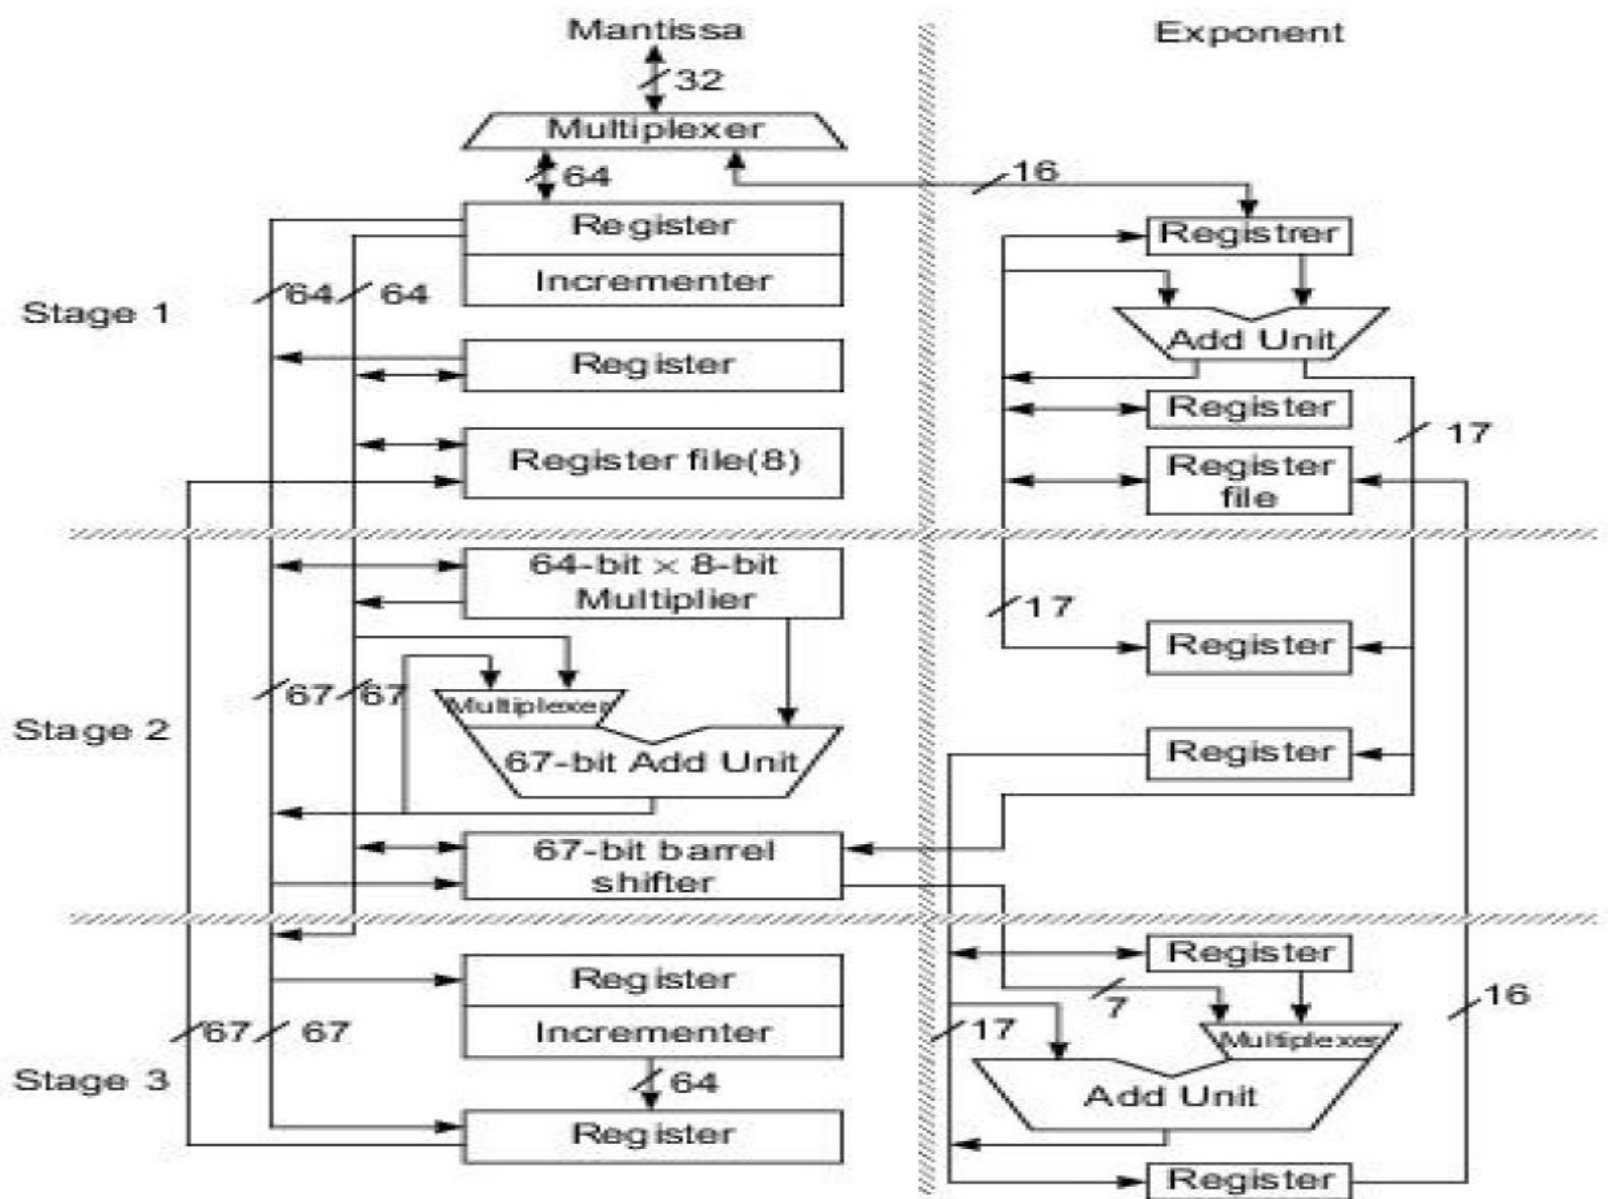
CSA = Carry save adder

CPA = Carry Propagate adder

$$P = A \times B$$

FLOATING POINT UNIT IN MOTOROLA MC68040

- This arithmetic pipeline has three stages.
- The mantissa section and exponent section are implemented as 2 separate pipelines
- The mantissa section can perform floating-point add or multiply operations on 32 bit or 64 bit



MANTISSA SECTION

- Stage 1
 - receives input operands and returns with computation results
 - 64-bit registers are used in this stage.
 - All three stages are connected using two 64-bit data buses.
- Stage 2
 - contains the array multiplier which is repeatedly used to carry out a long multiplication of the two mantissas.
 - The 67-bit adder performs the addition/subtraction of two mantissas
 - Barrel shifter is used for normalization.
- Stage 3
 - contains registers for holding results before they are loaded into the register file in stage 1 for subsequent use by other instructions.

EXPONENT SECTION

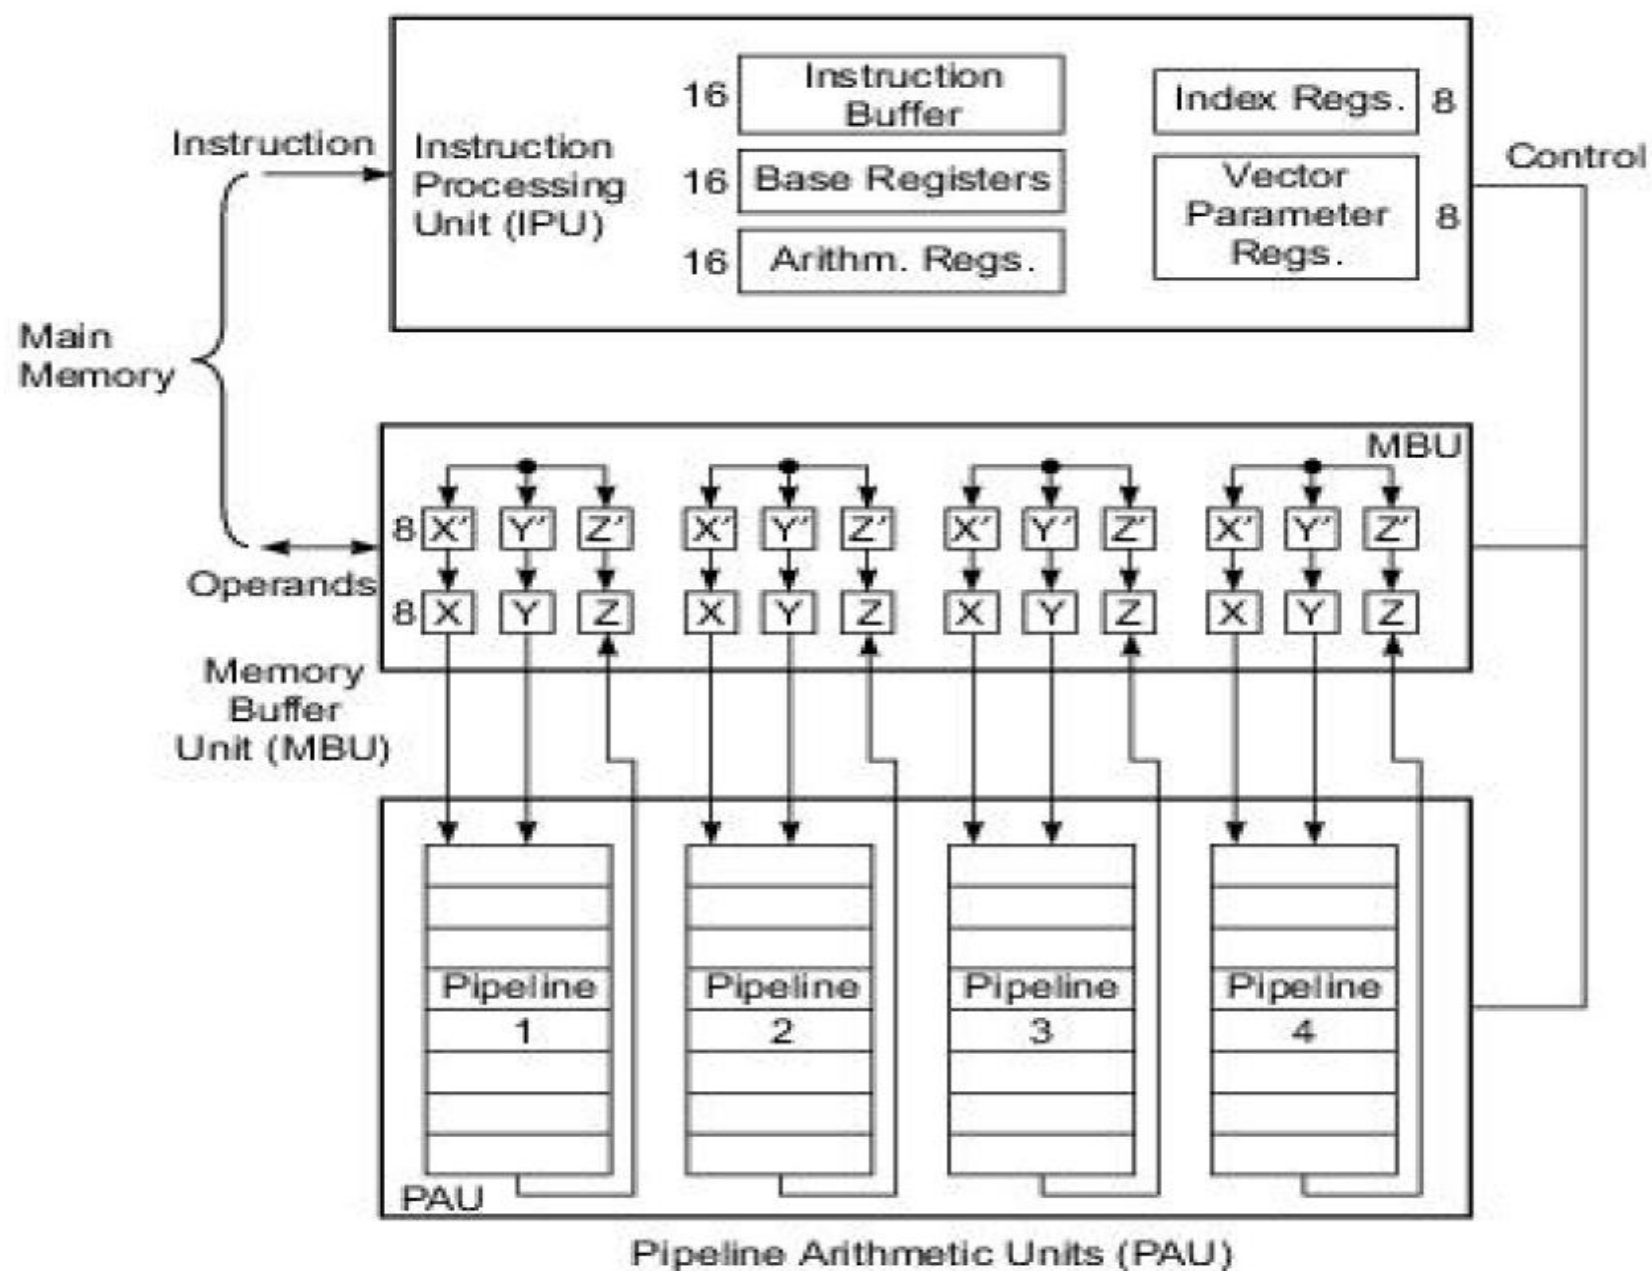
- 16 bit bus is used between stages.
- Stage 1
 - It has an exponent adder for comparing the relative magnitude of two exponents.
- Stage 2
 - The result of stage 1 is used to equalize the exponents before mantissa addition can be performed
 - Therefore, output of the exponent adder is sent to the barrel shifter for mantissa alignment.
- Stage 3
 - After normalization of the final result the exponent is readjusted in stage 3 using another adder.
 - The final value of the resulting exponent is fed from the register in stage 3 to the register file in stage 1

MULTIFUNCTIONAL **ARITHMETIC** PIPELINE

- A pipeline which can perform more than one function, is called as multifunctional pipeline
- It can be of 2 types
 - Static pipeline
 - Dynamic pipeline
- Static pipelines perform one function at a time
 - Different functions can be performed at different times.
- A dynamic pipeline allows several functions to be performed simultaneously through the pipeline
 - as long as there are no conflicts in the shared usage of pipeline stages

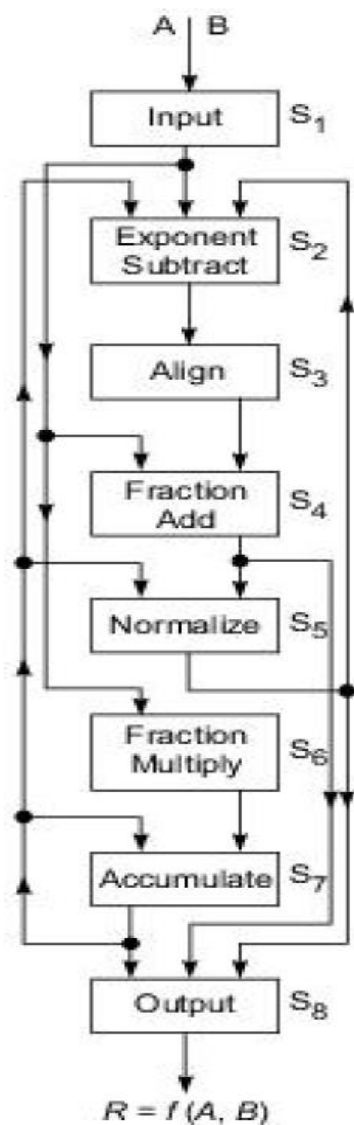
TI/ASC ARITHMETIC PROCESSOR DESIGN

- There are 4 pipeline units in this s/m
- Instruction processing unit
 - handled the fetching and decoding of instructions.
 - working registers in the processor controlled the operations of the memory buffer unit and arithmetic units.
- Operand buffers
 - 2 set of operand buffers in each arithmetic unit
 - {X,Y,Z}
 - {X',Y',Z'}
 - X',X, Y' and Y are used for input operands
 - Z' and Z were used to o/p results
 - intermediate results could be routed from Z-registers to either X or Y registers
- Both processor and memory buffers accessed the main memory for instructions and operands/results




○ Pipeline arithmetic unit

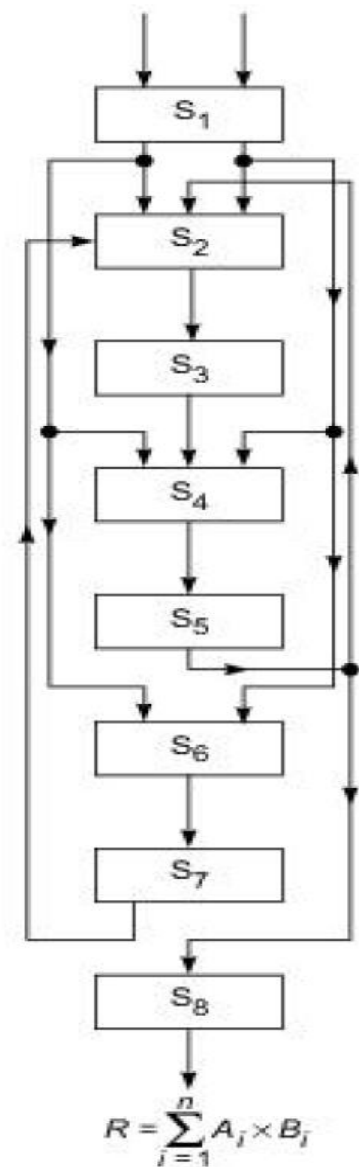
- Each pipeline arithmetic unit had eight stages
- PAU was a static multifunction pipeline which could perform only one function at a time.
- Both fixed-point and floating-point arithmetic functions could be performed by this pipeline.
- The PAU also supported vector in addition to scalar arithmetic operations.
- Different functions required different pipeline stages and different interstage connection patterns



(a) Pipeline stages and interconnections



(b) Fixed-point multiplication



(c) Floating-point dot product

EXAMPLE

- fixed-point multiplication required the use of only segments S1, S6, S7 and S8
- The floating-point dot product function, which performs the dot product operation between two vectors, required the use of all segments with the complex connections
 - This dot product was implemented by essentially the following accumulated summation of a sequence of multiplications through the pipeline